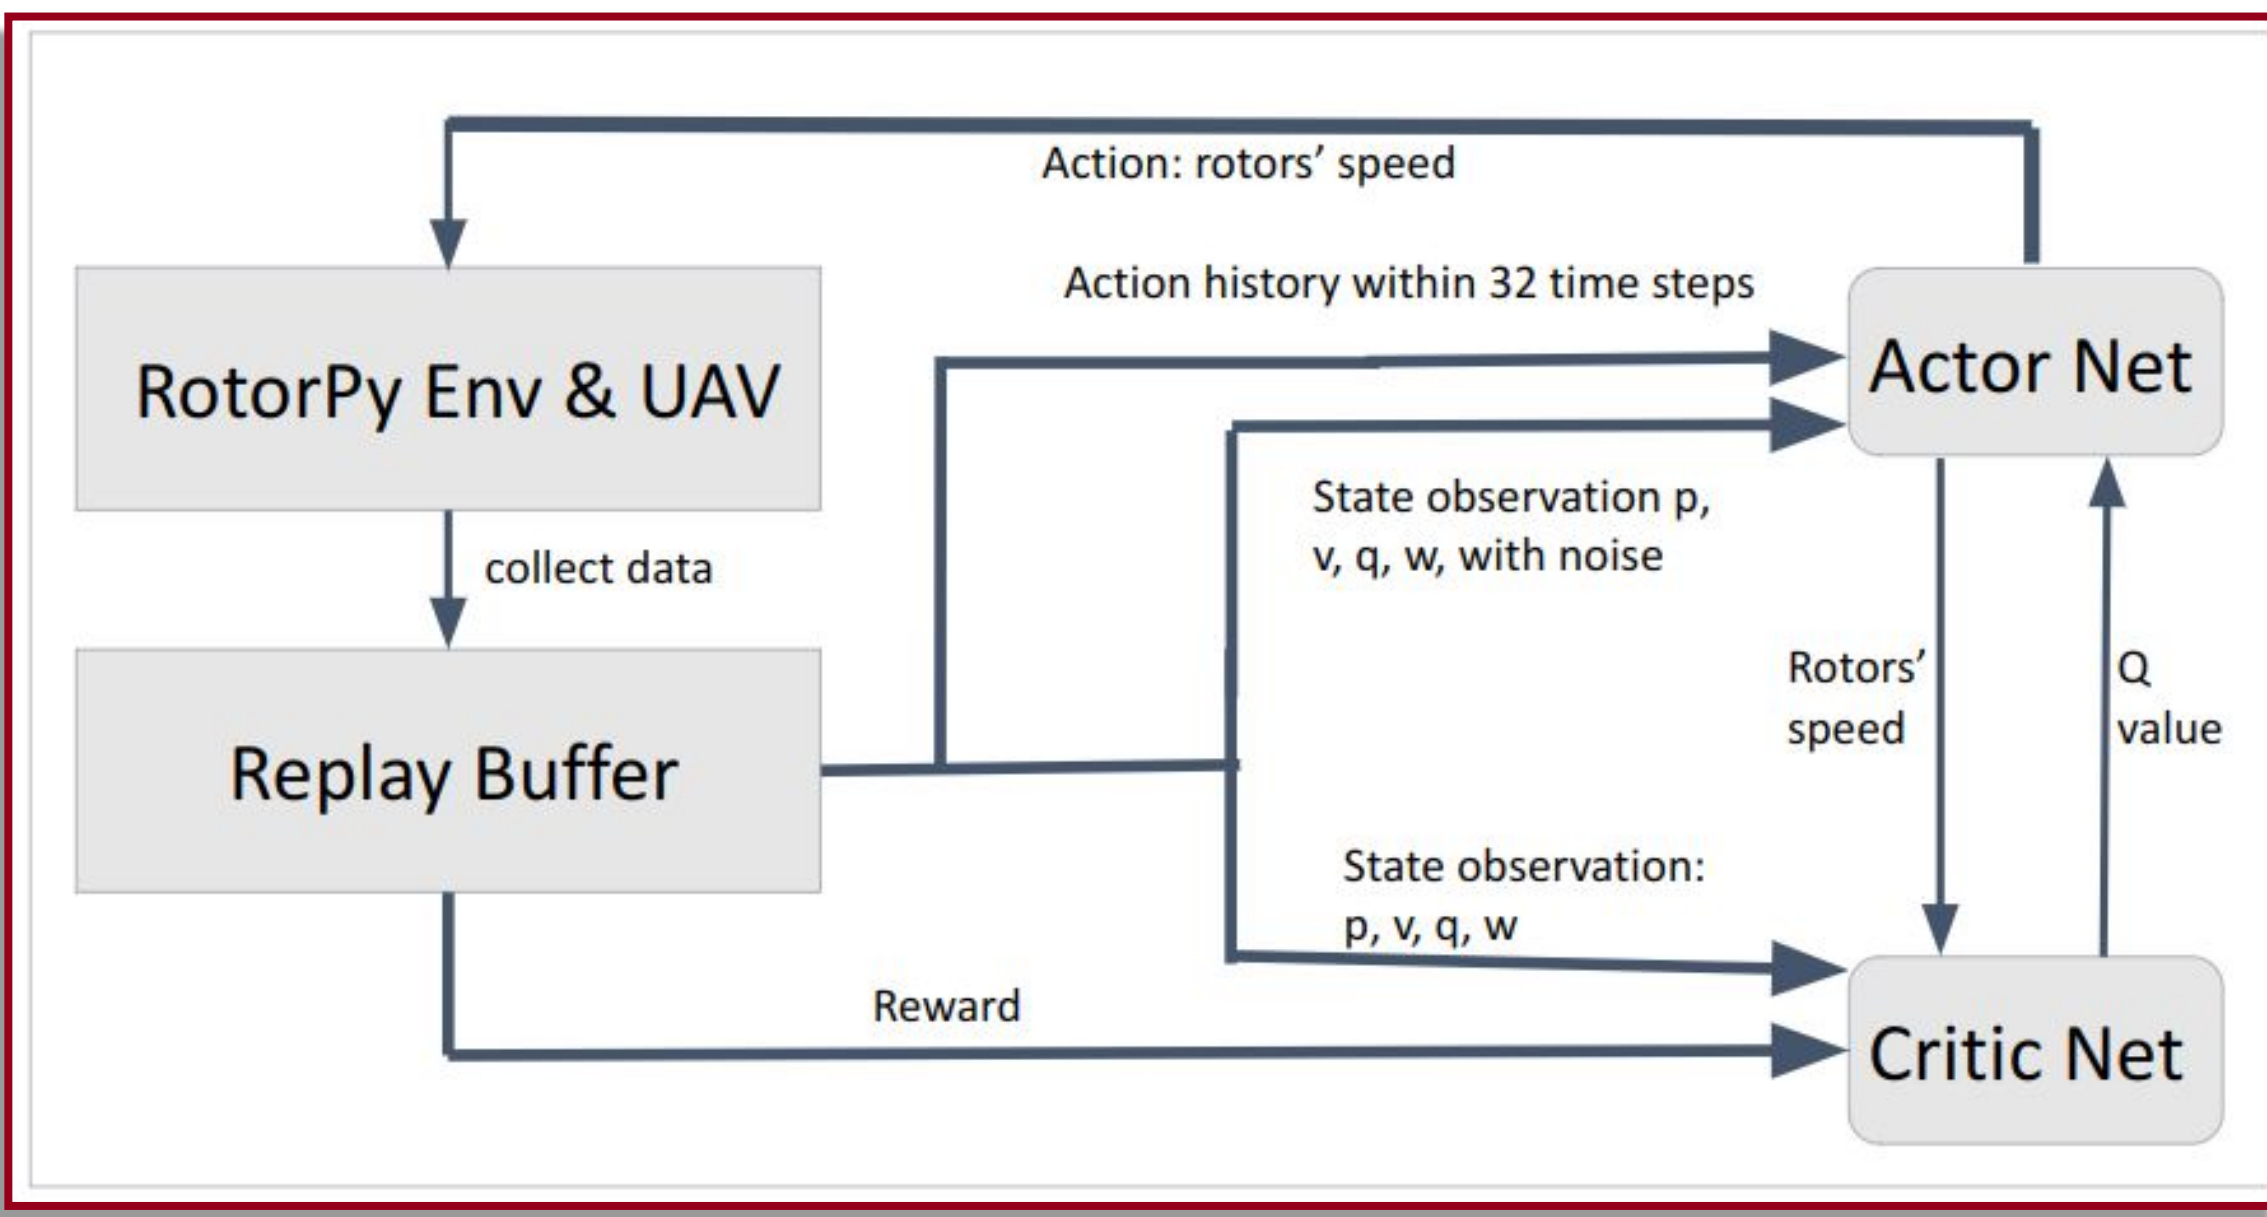


Learning Robust UAV Hovering in Wind-Disturbed Environments

Songhao Huang, Bowen Yang

INTRODUCTION/MOTIVATION

Hovering is essential for UAV applications like delivery and disaster response, where stability in unpredictable conditions is crucial. This challenge involves the complexity of end-to-end learning, adapting to varying winds, and stability across diverse initial conditions. We used a reinforcement learning approach to train an end-to-end policy for UAV hovering under wind disturbances. Compared to traditional methods, our policy achieves robust performance across varying winds and initial conditions.



METHODS

1. Simulation Environment

We used RotorPy [1], a Python-based multirotor simulation with aerodynamic wrenches. The state **observation** includes position, velocity, quaternion, and angular velocity, while the control **action** is defined as four rotor speeds.

We defined the **reward** function as the weighted sum of: current position error, current velocity error, the error of constant in quaternion, survival

METHODS Continued

reward, and goal reaching reward.

2. Architecture

Inspired by [2][3], we utilized an off-policy RL algorithm (TD3) with Actor and Critic networks, each comprising a 3-layer structure with 64 neurons per hidden layer.

3. Training Details

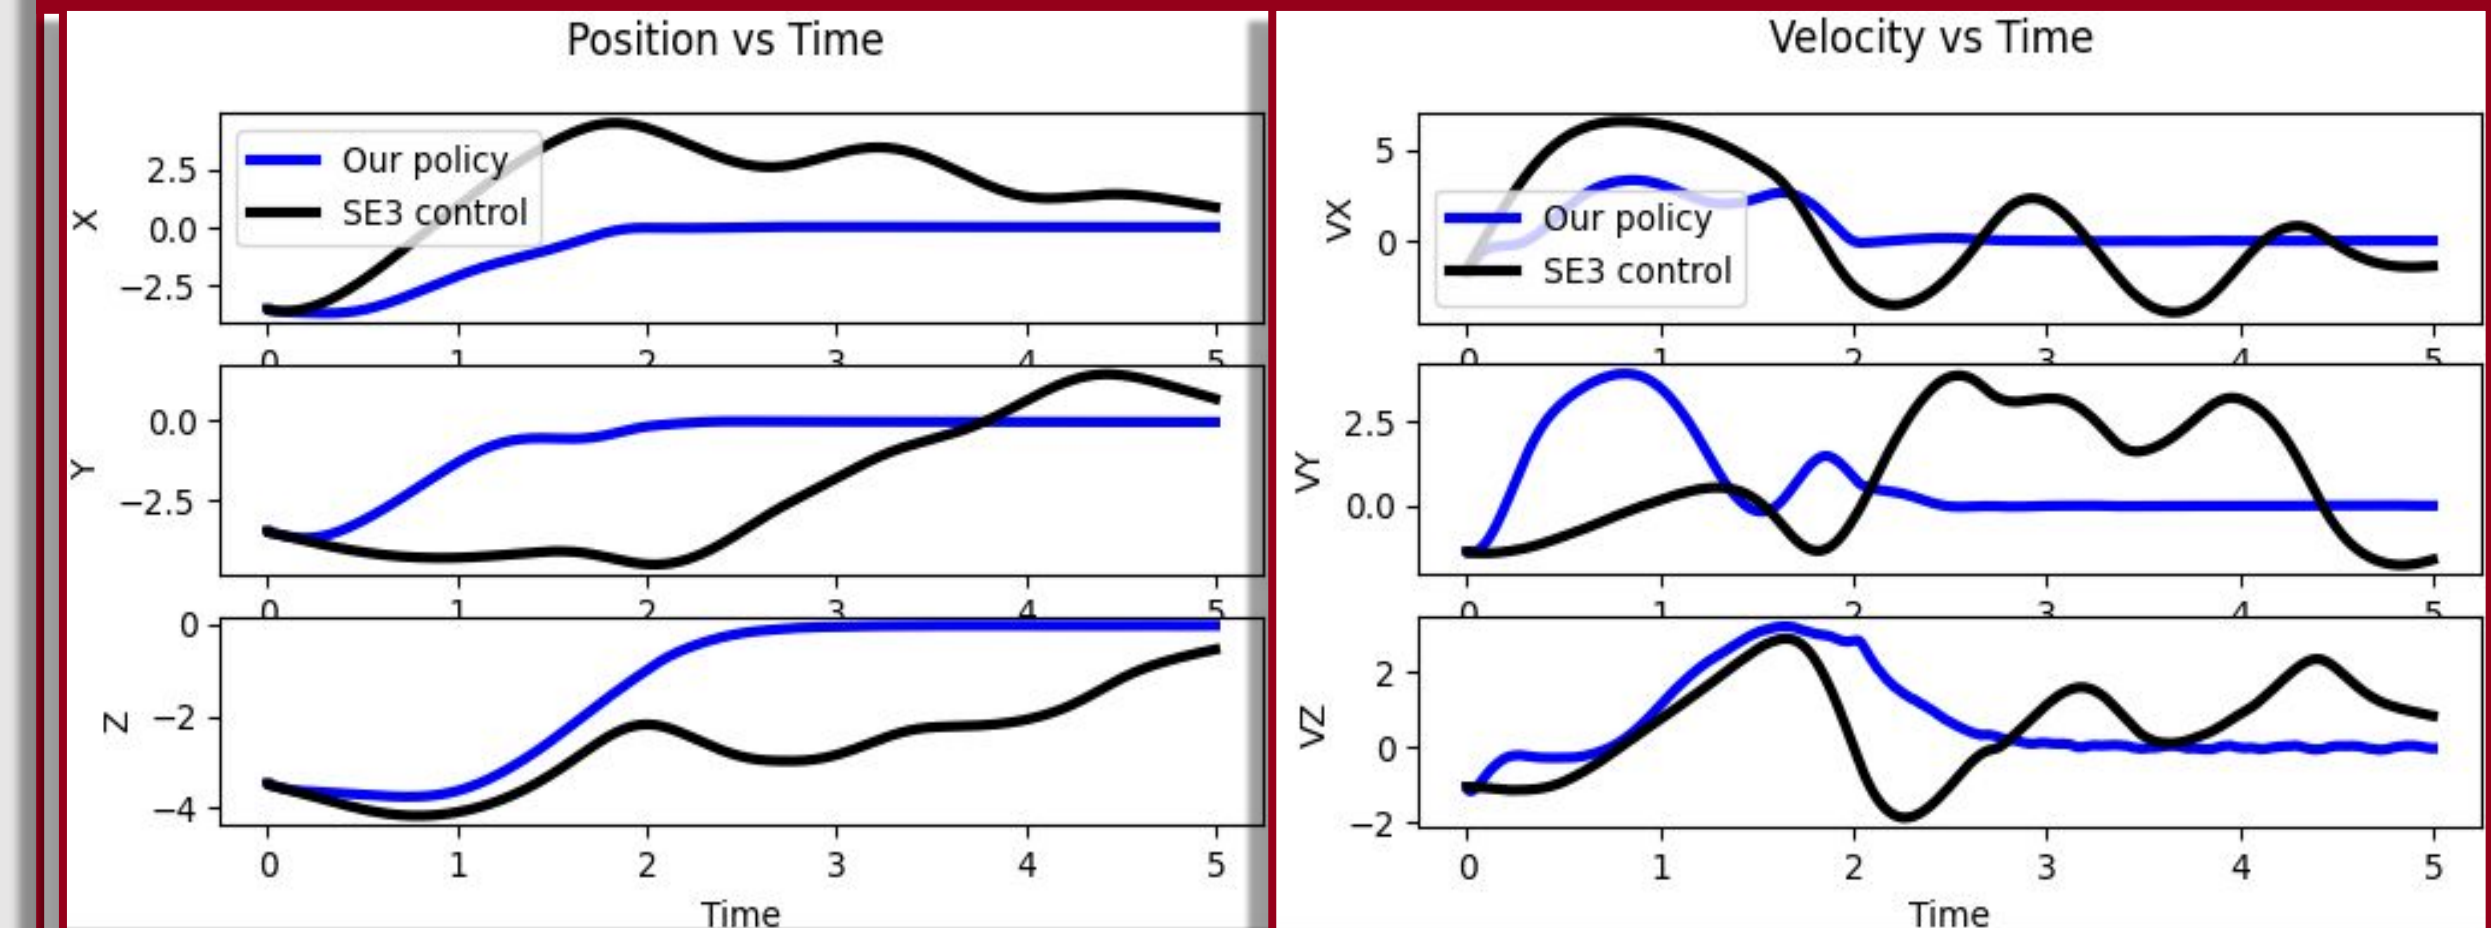
3.1 Guidance: We randomly initialize all the UAV state within a reasonable range, occasionally starting it in a hover position to enhance training.

3.2 Random noise is added to observations to improve policy robustness against wind variability.

3.3 Action history is fed into the policy to accelerate motor response, modeled as a first-order low-pass filter.

3.4 Curriculum learning adjusts reward weights and guidance probabilities to optimize training progression.

RESULTS Continued



Constant wind	Mild			Medium			Fierce		
	Success rate	Reaching time	Control cost	success rate	Reaching time	control cost	success rate	Reaching time	Control cost
Ours	100 %	2.57 s	222.99	100 %	3.01 s	268.02	57 %	2.99 s	249.66
SE3 ctrl	100 %	3.29 s	307.92	21 %	3.38 s	305.01	2 %	3.9 s	313.34

Sine wind	Mild			Medium			Fierce		
	Success rate	Reaching time	Control cost	success rate	Reaching time	control cost	success rate	Reaching time	Control cost
Ours	100 %	2.59 s	223.37	98 %	2.99 s	256.85	73 %	3.65 s	310.62
SE3 ctrl	100 %	3.25 s	302.49	79 %	3.66 s	341.29	21 %	4.1 s	382

Initial Condition	Ideal			Challenge			Adverse		
	Success rate	Reaching time	Control cost	success rate	Reaching time	control cost	success rate	Reaching time	Control cost
Ours	100 %	1.67 s	133.58	100 %	2.52 s	199.61	100 %	2.77 s	230.32
SE3 ctrl	100 %	2.16 s	163.19	74 %	3.27 s	297.59	50 %	3.18 s	285.64

4. Flaw and Future works

We also evaluated our policy on a path-tracking task by feeding offset position and velocity into the hover policy. It struggled with high-speed or steep curves. We tried to train a path-tracking policy using the RL framework and proved difficult due to poor convergence and generalization. Future works could involve fine-tuning the hover policy for path tracking or learning a residual policy to compensate for disturbances on a nominal controller's output.

5. Video, Code and References:

github: https://github.com/KANZEZ/rl_uav_ctrl

References links:

[1] <https://arxiv.org/abs/2306.04485>

[2] <https://arxiv.org/pdf/2311.13081>

[3] <https://arxiv.org/pdf/1802.09477>

[4] <https://ieeexplore.ieee.org/document/5717652>

Youtube



RESULTS

1. Training Time

Our policy was trained on a laptop with a RTX4060 GPU for **40 mins** (~700000 iteration)

2. Simulation Benchmark

We compare the performance with SE3 controller[4]. The following are corresponding position/velocity vs. time plots in **one test**, with the same adverse random initial conditions and constant wind disturbance.

3. More Experiments

The table shows the performance for two controllers testing for **100 times** in each case.